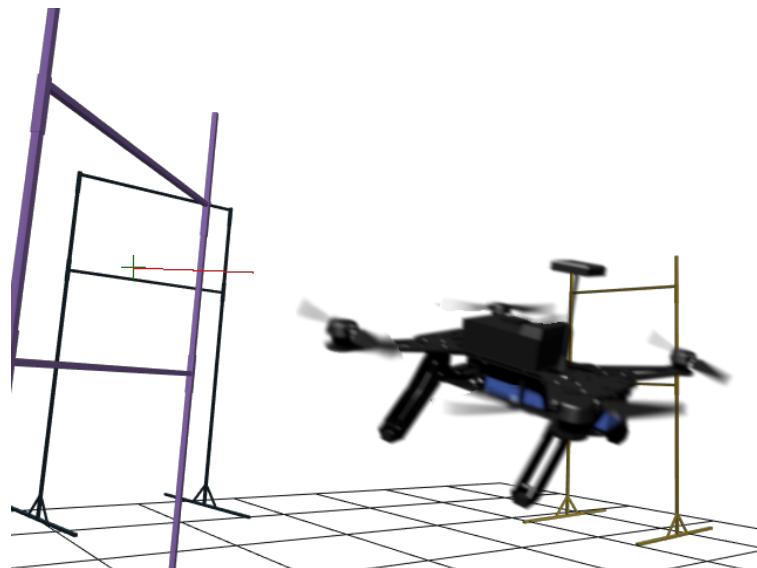




AARHUS
UNIVERSITY

DEPARTMENT OF ENGINEERING



Generating synthetic images to train convolutional neural networks for autonomous drone racing

Transferring knowledge from a semi-synthetic domain into the real world for trajectory planning

Master's thesis in M.Sc. Computer Engineering

Théo Morales

MASTER'S THESIS 2019:NN

Generating synthetic images to train convolutional neural networks for autonomous drone racing

Transferring knowledge from a semi-synthetic domain into the real world for
trajectory planning

Théo Morales



Department of Engineering
Control and Automation Laboratory
AARHUS UNIVERSITY
Aarhus, Denmark - 2019

Generating synthetic images to train Convolutional Neural Networks for autonomous drone racing

Transferring knowledge from a semi-synthetic domain into the real world for drone steering
Théo Morales

© Théo Morales, 2019.

Supervisor: Erdal Kayacan, Department of Engineering

Co-Supervisor: Andriy Sarabakha, Department of Engineering

External Examiner: Name, Department

Master's Thesis 2019:NN

Department of Engineering

Control and Automation Laboratory

Aarhus University

DK-8000 Aarhus

Telephone +4541893252

Cover: Vision-based gate detection for autonomous drone racing.

Typeset in L^AT_EX

Aarhus, Denmark 2019

Using synthetic images to train convolutional neural networks for autonomous drone racing
Transferring knowledge from a semi-synthetic domain into the real world for drone steering
Theo Morales

Department of Engineering
Aarhus University

Abstract

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Keywords: lorem, ipsum, dolor, sit, amet, consectetur, adipisicing, elit, sed, do.

Acknowledgements

 Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Theo Morales, Aarhus, June 2019

Contents

List of Figures	VII
List of Tables	IX
1 Introduction	1
2 Autonomous Drone Racing	3
2.1 What is drone racing?	3
2.1.1 Stakes	4
2.1.2 Challenges	4
2.1.3 Beyond the hype and into reality	4
2.2 Litterature review	4
3 Convolutional Neural Networks	5
3.1 Advantages and applications	5
3.2 Architecture	5
3.2.1 Typical inputs	5
3.2.2 Typical outputs	5
3.2.3 Layer types	5
3.2.4 Regularization	5
3.3 Training and evaluation	5
3.3.1 Dataset choice	5
3.3.2 Hyperparameter tuning	5
3.4 Common pitfalls	5
3.4.1 Underfitting	5
3.4.2 Overfitting	5
4 Choice of topology	6
4.1 Pretrained models	6
4.2 Custom models	6
4.3 DroNet: bakdwjwkdkw	6
4.4 Modifying DroNet	6
5 Dataset generation	7
5.1 Problem statement	7
5.2 Proposed solution	7
5.2.1 Advantages	7
5.2.2 Limitations	7
5.3 Synthetic gates generation	7
5.3.1 OpenGL scene	7
5.3.2 Modelling the camera	7

5.3.3	Mapping a virtual space	7
5.4	Foreground blending	7
6	Network training	8
6.1	Dataset size and network capacity	8
6.2	Choice of Hyperparameters	8
6.3	Overcoming the overfitting problem	8
6.4	Achieving generalization	8
6.5	Evaluating on real world data	8
7	Control and trajectory planning	9
7.1	Problem statement	9
7.2	Robots Operating System	9
7.3	Proportional Integral Derivative controllers	9
7.4	Putting things together	9
8	Experiments and Discussion of results	10
8.1	Gate detection	10
8.2	Trajectory planning	10
9	Conclusion	11
Bibliography		XI
A	Appendix 1	XI

List of Figures

1.1	Drone racing arena at IROS 2016.	1
1.2	The Sense, Plan, Act robotic paradigm.	1
2.1	Drone Racing League arena in London (Steven Paston/PA)	3
2.2	The AiR Lab and the custom built obstacles around the author in Skejby	4

List of Tables

List of Tables

1

Introduction

Each year, the robotic community gathers at conferences such as IROS (International Conference on Intelligent Robots and Systems), where they exhibit their discoveries and achievements in the field. For the year 2019, the Drone Racing League will be hosting a series of races which autonomous drones will compete and attempt to outperform a professional drone pilot.

The company Lockheed Martin is sponsoring the event and offering a prize of 2 million USD



Figure 1.1: Drone racing arena at IROS 2016.

for the winning team. Teams of university students and other drone enthusiasts shall present innovative approaches on vision-based systems for unmanned aerial vehicles: hence, they showcase their progress through racing.

Most robotic system's operation can be modeled by the commonly known *Sense, Plan, Act* paradigm:

Talk about the predictions for drone racing and drones in general for the future

Add custom act, sense, plan paradigm pic

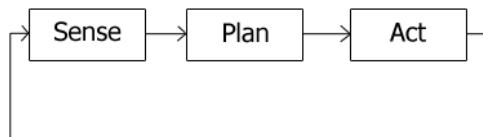


Figure 1.2: The Sense, Plan, Act robotic paradigm.

- **Sense:** gather information using sensors (camera, IMU, sonar...).
- **Plan:** create a world model using all the information, and plan the next move.
- **Act:** carry out the actions that the plan calls for.

This thesis will be focused mainly on the sensing of the drone control, and more precisely using computer vision algorithms along with machine learning concepts. However, an important part

of the trajectory planning directly follows the sensing phase, therefore those two first phases of the control loop can be included in the scope of the project.

To this day, most drones in the robotic research field are running *Robot Operating System* (ROS), which is a set of "libraries and tools to help software developers create robot applications". The implementation of the computer vision algorithms will be constrained within ROS, and cooperate with the different control components provided by the project team.

As Deep Learning has known an exponential growth during the past decade, computer vision applications tend to exploit the power of convolutional neural networks more and more. Thanks to their impressive accuracy in specific problem solving, CNNs are becoming the main choice for tasks such as: object detection, object segmentation, object recognition, etc...

In drone racing, the latest works, which are often the best performing, tend to employ CNNs for the sensing part, or even as an all-in-one solution (refer to).

The solution will be implemented iteratively, by gradually increasing the complexity of the vision algorithms, such that progress can be made even if the end goal is too difficult to achieve.

The project should attempt to fulfill the following goals:

- Recognize gates in the input image
- Detect and localize the closest gate's center
- Evaluate the closest gate's orientation and distance
- Plan a trajectory for the drone to fly across the gate center
- Refine the trajectory in real time, for an optimal flight time
- Make the drone follow that trajectory while adjusting in real time

However, the following functionalities are out of the scope and will be discarded:

- Detect any other obstacles in the input image
- Avoid obstacles that are on the drone's trajectory
- Map and localize the detected gates in world frame
- Plan a trajectory of a sequence of gates
- Apply online learning for adaptive racing

The focus of this research will be on developing a robust and generalized vision algorithm for detecting obstacles in a configurable manner and regardless of their shape and color, such that the race gates can be specified prior to the training of the algorithm, and thus tailored to the challenge itself. The end goal is to greatly facilitate the learning of the vision-based detection model, by automating the tedious and long process of collecting and annotating a large dataset required to train the object detection algorithm using Deep Learning.

Indeed, this task is often the most time-consuming part of training convolutional neural networks, mostly because each image must be annotated to be used in supervised learning. By removing this weight off researchers' shoulders, much more time is left to focus and experiment on the *planning* and *acting* problems of the robotic paradigm .

In this work, the efficiency and robustness of tailored dataset generation for Deep Learning-based computer vision algorithms, for the specific case of autonomous drone racing, is analysed and discussed. A simple approach to solving the *planning* challenge is also undertaken, which supports the praises of the proposed solution.

2

Autonomous Drone Racing

2.1 What is drone racing?

As in any race, competitors measure their skills at maneuvering their vehicle, from within or at distance, against the clock or against themselves, in an attempt to complete a circuit before everyone else.

In drone racing, the circuit is comprised of a set of obstacles, static or dynamic, usually arranged in a loop. For the case of classic drone racing, huge futuristic-looking arenas are used as a playground for extremely skilled pilots, who make the race look impressive and dynamic by the use of FPV goggles to fly the drones at very high speed, reaching 120mph [<https://thedroneracingleague.com/learn-more/>].

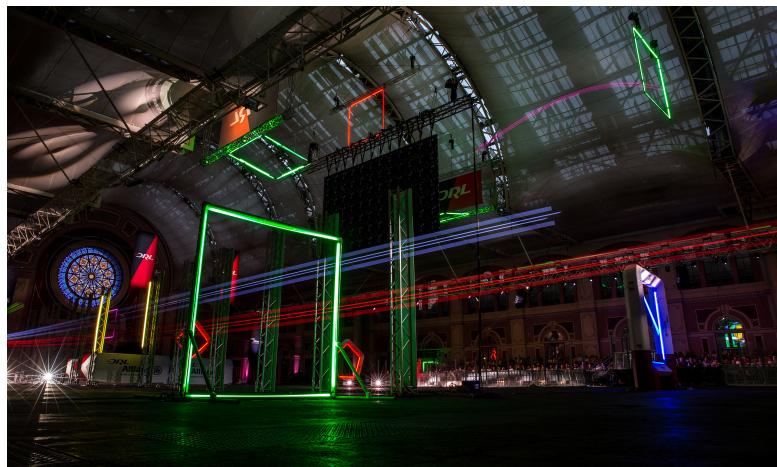


Figure 2.1: Drone Racing League arena in London (Steven Paston/PA)

For the case of autonomous drone racing, up until the *Drone Racing League* announced their autonomous competition of 2019, arenas looked much more like a laboratory than a glowing circuit, as it can be seen on Figure 2.2 showing the AiR Lab in Skejby (Aarhus, DK), with its gates built and painted for this work. These obstacles were inspired of the previous competitions organized by IROS, as seen on Figure 1.1 which also shows the typical type of drones used in this research area: consumer drones for the larger audience. This kind of drones is not made for sportive flight, but mostly for an easy hands-on experience, providing a relatively slow but stable flight.

The goal is to complete a given amount of turns as fast as possible, but since it is still the early days of drone racing and the costs of a potential crash would be too high, each drone races against the clock, one after another. Ultimately, autonomous drones will be racing against skilled pilots, as the *Drone Racing League* seems to encourage it by giving a total prize of 2 millions USD for the race of 2019.

cite image:
<https://homegadgets/tech-news/the-drone-racing-league-has-set-a-new-world-record-for-the-fastest-drone-around-113641968811>



Figure 2.2: The AiR Lab and the custom built obstacles around the author in Skejby

2.1.1 Stakes

2.1.2 Challenges

2.1.3 Beyond the hype and into reality

2.2 Litterature review

3

Convolutional Neural Networks

3.1 Advantages and applications

3.2 Architecture

3.2.1 Typical inputs

3.2.2 Typical outputs

3.2.3 Layer types

3.2.4 Regularization

3.3 Training and evaluation

3.3.1 Dataset choice

3.3.2 Hyperparameter tuning

3.4 Common pitfalls

3.4.1 Underfitting

3.4.2 Overfitting

4

Choice of topology

4.1 Pretrained models

4.2 Custom models

4.3 DroNet: bakdwjwkdkw

4.4 Modifying DroNet

5

Dataset generation

5.1 Problem statement

5.2 Proposed solution

5.2.1 Advantages

5.2.2 Limitations

5.3 Synthetic gates generation

5.3.1 OpenGL scene

5.3.2 Modelling the camera

5.3.3 Mapping a virtual space

5.4 Foreground blending

6

Network training

- 6.1 Dataset size and network capacity
- 6.2 Choice of Hyperparameters
- 6.3 Overcoming the overfitting problem
- 6.4 Achieving generalization
- 6.5 Evaluating on real world data

7

Control and trajectory planning

7.1 Problem statement

7.2 Robots Operating System

7.3 Proportional Integral Derivative controllers

7.4 Putting things together

8

Experiments and Discussion of results

8.1 Gate detection

8.2 Trajectory planning

9

Conclusion

A

Appendix 1